

ADVANCED UNIX PROGRAMMING MIDTERM REPORT



TEAM 9 — 林禾堃、馬毓昇、陳曦

Dec 2023

1. Question 1

The code use `pthread_once` function to ensure that the `pthread` key will only create once. After the creation of key, `envbuf` should be `NULL`, and therefore `malloc` a space, associated with the key we've just created.

Then every time we run the `getenv` function, we only write the result on thread-specific data `envbuf`. Then we can ensure that even if two threads call this function at the same time, since the result is written in thread-specific data, there won't be inconsistencies.

2. Question 2

No, the reason why this version of `getenv` is still not `async-signal` safety is because `malloc` function is not `async-signal` safe. For example, when two threads run `malloc` at the same time, there might be race condition, and therefore not `async-signal` safe. Therefore simply block signals won't solve this issue.

3. Question 3

The program cannot run successfully on FreeBSD, to understand why, we tried to use `gdb-peda` to trace and see what happened while running the program.

We starts with checking function symbols. Below is the result, we found that there are multiple definitions of `getenv` function.

For example, in both `stdlib` and `rtld-elf` we found the definition, and also our implementation.

```
gdb-peda$ info functions ^getenv
All functions matching regular expression "^getenv":

File /usr/src/lib/libc/stdlib/getenv.c:
419:   char *getenv(const char *);

File /usr/src/libexec/rtld-elf/rtld.c:
6320:  char *getenv(const char *);

Non-debugging symbols:
0x00000000400da4   getenv
0x00000000404ac300  getenv@plt
0x000000004068ec00  getenv@plt
```

Then we figure out that even we don't send any arguments, which suppose to print out help message then return directly, but it stuck rather than terminating.

So we set a break point at getenv function, then run debugger. And we found that the program keep stuck in getenv function as below.

```
gdb-peda$ c
Continuing.

Breakpoint 1.1, 0x000000000400da8 in getenv ()
gdb-peda$ c
Continuing.

Breakpoint 1.1, 0x000000000400da8 in getenv ()
gdb-peda$ c
Continuing.

Breakpoint 1.1, 0x000000000400da8 in getenv ()
gdb-peda$ c
Continuing.

Breakpoint 1.1, 0x000000000400da8 in getenv ()
gdb-peda$ c
Continuing.
```

To see more details, use ni to see next instruction. Then we found that the program will keep running between 0x400da8 to 0x400dd0.

```
Breakpoint 1.1, 0x000000000400da8 in getenv ()
gdb-peda$ ni
0x000000000400dac in getenv ()
gdb-peda$
0x000000000400db0 in getenv ()
gdb-peda$
0x000000000400db4 in getenv ()
gdb-peda$
0x000000000400db8 in getenv ()
gdb-peda$
0x000000000400dbc in getenv ()
gdb-peda$
0x000000000400dc0 in getenv ()
gdb-peda$
0x000000000400dc4 in getenv ()
gdb-peda$
0x000000000400dc8 in getenv ()
gdb-peda$
0x000000000400dcc in getenv ()
gdb-peda$
0x000000000400dd0 in getenv ()
gdb-peda$
Breakpoint 1.1, 0x000000000400da8 in getenv ()
```

Below are the instructions between these two addresses.


```

gdb-peda$ x/67i 0x400da4
0x400da4 <getenv>: stp x29, x30, [sp, #-64]!
=> 0x400da8 <getenv+4>: adrp x1, 0x400000
0x400dac <getenv+8>: add x1, x1, #0xd90
0x400db0 <getenv+12>: mov x29, sp
0x400db4 <getenv+16>: stp x19, x20, [sp, #16]
0x400db8 <getenv+20>: adrp x19, 0x411000
0x400dbc <getenv+24>: stp x23, x24, [sp, #48]
0x400dc0 <getenv+28>: add x23, x19, #0x340
0x400dc4 <getenv+32>: mov x24, x0
0x400dc8 <getenv+36>: add x0, x23, #0x8
0x400dcc <getenv+40>: stp x21, x22, [sp, #32]
0x400dd0 <getenv+44>: bl 0x400940 <pthread_once@plt>
0x400dd4 <getenv+48>: add x0, x23, #0x18

```

Also, when back tracing, we see the same situation.

```

gdb-peda$ bt
#0 __thr_interpose_libc () at /usr/src/lib/libthr/thread/thr_syscalls.c:643
#1 0x000000004049ec40 in _libpthread_init (curthread=0x0) at /usr/src/lib/libthr/thread/thr_init.c:329
#2 0x00000000404a4204 in _thr_check_init () at /usr/src/lib/libthr/thread/thr_private.h:930
#3 _thr_once (once_control=0x411348 <init_done>, init_routine=0x400d90 <thread_init>) at /usr/src/lib/libthr/thread/thr_once.c:72
#4 0x000000004040dd4 in getenv ()
#5 0x000000004049ed24 in init_private () at /usr/src/lib/libthr/thread/thr_init.c:501
#6 _libpthread_init (curthread=0x0) at /usr/src/lib/libthr/thread/thr_init.c:332
#7 0x00000000404a4204 in _thr_check_init () at /usr/src/lib/libthr/thread/thr_private.h:930
#8 _thr_once (once_control=0x411348 <init_done>, init_routine=0x400d90 <thread_init>) at /usr/src/lib/libthr/thread/thr_once.c:72
#9 0x000000004040dd4 in getenv ()
#10 0x000000004049ed24 in init_private () at /usr/src/lib/libthr/thread/thr_init.c:501
#11 _libpthread_init (curthread=0x0) at /usr/src/lib/libthr/thread/thr_init.c:332
#12 0x00000000404a4204 in _thr_check_init () at /usr/src/lib/libthr/thread/thr_private.h:930
#13 _thr_once (once_control=0x411348 <init_done>, init_routine=0x400d90 <thread_init>) at /usr/src/lib/libthr/thread/thr_once.c:72
#14 0x000000004040dd4 in getenv ()

```

The reason why this issue occur is that some initialize function tried to get environment variable by calling getenv function. But in our implementation, we also use some functions that need initialize by calling getenv function (for example, pthread functions and malloc.), and therefore a dead loop occurs.

We can verify that by output the name in getenv function.

```

root@genet0:~/Advanced-UNIX-Programming/HW10 # ./assignment10
MALLOC_CONF
LIBPTHREAD_BIGSTACK_MAIN
LIBPTHREAD_SPLITSTACK_MAIN
LIBPTHREAD_SPINLOOPS
LIBPTHREAD_YIELDLOOPS
LIBPTHREAD_QUEUE_FIFO

```

So in order to solve this issue, we should make sure that when the initialize functions calling getenv, we won't run thread-safe codes, while after that, just do thread-safe version.

First, declare a flag function in global.

```
// for preventing initialize process dead
static int init_flag = 1;
```

Then since we cannot run malloc, so instead we allocate a static char array, so that even the function return, the value still holds.

```
// buffer for initialize process
static char buf[MAXSTRINGSZ];
```

For initialize functions, just assign the buf to envbuf, otherwise we run the thread-safe code.

```
if(init_flag == 0){
    pthread_once(&init_done, thread_init);
    pthread_mutex_lock(&env_mutex);
    envbuf = (char *)pthread_getspecific(key);

    if (envbuf == NULL) {
        envbuf = malloc(MAXSTRINGSZ);
        if (envbuf == NULL) {
            pthread_mutex_unlock(&env_mutex);
            return(NULL);
        }
        pthread_setspecific(key, envbuf);
    }
}
else{
    envbuf = buf;
}
```

Whenever we are going to run functions like pthread, check whether running initialization or not.

```
len = strlen(name);
for (i = 0; environ[i] != NULL; i++){
    if ((strcmp(name, environ[i], len) == 0) &&
        (environ[i][len] == '=')) {
        strncpy(envbuf, &environ[i][len + 1], MAXSTRINGSZ - 1);
        if(init_flag == 0){
            pthread_mutex_unlock(&env_mutex);
        }
        return(envbuf);
    }
}

if(init_flag == 0){
    pthread_mutex_unlock(&env_mutex);
}

return(NULL);
}
```

When running to main function, we shall assume that all initialization are done, so set the initialization flag to 0.

```
// set init flag
init_flag = 0;
```

Below is the output with debug message.

```
root@genet0:~/Advanced-UNIX-Programming/HW10 # ./assignment10 PATH
MALLOC_CONF
LIBPTHREAD_BIGSTACK_MAIN
LIBPTHREAD_SPLITSTACK_MAIN
LIBPTHREAD_SPINLOOPS
LIBPTHREAD_YIELDLOOPS
LIBPTHREAD_QUEUE_FIFO
PATH = /sbin:/bin:/usr/sbin:/usr/bin:/usr/local/sbin:/usr/local/bin:/root/bin
```